

Serializability

Serializability

- When multiple transactions are running concurrently then there is a possibility that the database may be left in an inconsistent state.
- Serializability is a concept that helps us to check which schedules are serializable.
- A serializable schedule is the one that always leaves the database in consistent state.

Serializability

When multiple transactions are being executed by the operating system in a multiprogramming environment, there are possibilities that instructions of one transaction are interleaved with some other transaction

➤ **Schedule** – **A chronological (sequential) execution sequence of a transaction is called a schedule.** A schedule can have many transactions in it, each comprising of a number of instructions/tasks.

➤ **Serial Schedule** – It is **a schedule in which transactions are aligned in such a way that one transaction is executed first.** When the first transaction completes its cycle, then the next transaction is executed. Transactions are ordered one after the other. This type of schedule is called a serial schedule, as transactions are executed in a serial manner.

Concurrent Transactions

- In a multi-transaction environment, serial schedules are considered as a benchmark.
- The execution sequence of an instruction in a transaction cannot be changed, but two transactions can have their instructions executed in a random fashion.
- This execution does no harm if two transactions are mutually independent and working on different segments of data.
- But in case these two transactions are working on the same data, then the results may vary.

Checking Serializability

➤ Let $T1$ and $T2$ be two transactions that transfer funds from one account to another. Transaction $T1$ transfers \$50 from account A to account B . It is defined as:

$T1$: read(A);

$A := A - 50$;

write(A);

read(B);

$B := B + 50$;

write(B)

Transaction $T2$ transfers 10 percent of the balance from account A to account B . It is defined as:

$T2$: read(A);

$temp := A * 0.1$;

$A := A - temp$;

write(A);

read(B);

$B := B + temp$;

write(B)

Checking Serializability

➤ Suppose the current values of accounts A and B are \$1000 and \$2000, respectively. Suppose also that the two transactions are executed one at a time in the order T_1 followed by T_2 .

T_1	T_2
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code>

Figure 1 Schedule 1 – A serial schedule in which T_1 is followed by T_2

Checking Serializability

- If the transactions are executed one at a time in the order T2 followed by T1, then the corresponding execution sequence is that of Figure 2.

T_1	T_2
<pre>read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) commit</pre>	<pre>read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) commit</pre>

Figure 2 Schedule 2 – A serial schedule in which a T2 is followed by T1

Checking Serializability

- One possible schedule appears in Figure 3.

T_1	T_2
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B) commit	read(B) $B := B + temp$ write(B) commit

Figure 3 Schedule 3 – A concurrent schedule equivalent to schedule 1

Checking Serializability

➤ Consider the schedule of Figure 4

T_1	T_2
<pre> read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) commit </pre>	<pre> read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) commit </pre>

Figure 4 Schedule 4 – A concurrent schedule resulting in an inconsistent state

It is the job of the database system to ensure that any schedule that is executed will leave the database in a consistent state. The concurrency-control component of the database system carries out this task.

Conflict Serializability

- Serial schedules are serializable, but if steps of multiple transactions are interleaved, it is harder to determine whether a schedule is serializable.
- It is difficult to determine exactly what operations a transaction performs and how operations of various transactions interact.
- We will consider only two operations: read and write
- Between a **read(Q)** instruction and a **write(Q)** instruction on a data item **Q**, a transaction may perform an arbitrary sequence of operations on the copy of Q that is residing in the local buffer of the transaction

Conflict Serializability

- A schedule is called conflict serializable if it can be transformed into a serial schedule by swapping non-conflicting operations.
- **Conflicting operations:** Two operations are said to be conflicting if all conditions satisfy:
 - They belong to different transactions
 - They operate on the same data item
 - At Least one of them is a write operation

Conflict Equivalent Schedule

- A schedule is called conflict serializable if it can be transformed into a serial schedule by swapping non-conflicting operations.

S	
T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	

S ≡ S'

S'	
T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
	W(A)

Conflict Serializability

- Let us consider a schedule S in which there are two consecutive instructions, I and J , of transactions T_i and T_j , respectively ($i \neq j$)
- If I and J refer to different data items, then we can swap I and J without affecting the results of any instruction in the schedule
- if I and J refer to the same data item Q , then the order of the two steps may matter
- Since we are dealing with only read and write instructions, there are four cases that we need to consider:

Conflict Serializability

R(A)	R(A)	Non Conflicting
R(A)	W(A)	
W(A)	R(A)	
W(A)	W(A)	Conflicting
R(B)	W(A)	
W(B)	R(A)	
R(B)	W(A)	Non Conflicting
W(A)	W(B)	

Conflict Serializability

1. $I = \text{read}(Q)$, $J = \text{read}(Q)$. The order of I and J does not matter, (same value of Q is read by T_i and T_j)
2. $I = \text{read}(Q)$, $J = \text{write}(Q)$. If I comes before J , then T_i does not read the value of Q that is written by T_j in instruction J . If J comes before I , then T_i reads the value of Q that is written by T_j . (order of I and J matters)
3. $I = \text{write}(Q)$, $J = \text{read}(Q)$. (order of I and J matters)
4. $I = \text{write}(Q)$, $J = \text{write}(Q)$. (order doesn't affect T_i & T_j) However, the value obtained by the next $\text{read}(Q)$ instruction of S is affected, since the result of only the latter of the two write instructions is preserved in the database.

Conflict Serializability

S	
T1	T2
R(A) W(A)	R(A) W(A)
R(B)	

S'	
T1	T2
R(A) W(A) R(B)	R(A) W(A)

Check adjacent Non-conflicting pair

Conflict Serializability

- I and J **conflict** if :
 - ✓ They are operations belong to different transactions
 - ✓ Access the same data item
 - ✓ At least one of these instructions is a write operation
- Consider Schedule in the form of only read and write operations:

T_1	T_2
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B) commit	read(B) $B := B + temp$ write(B) commit

Schedule 1

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	read(B) write(B)

Schedule 2: Only read and write operations

Conflict Serializability

- Let I and J be consecutive instructions of a schedule S . If I and J are instructions of different transactions and I and J do not conflict, then we can swap the order of I and J to produce a new schedule S'
- S is equivalent to S' , since all instructions appear in the same order in both schedules except for I and J
- The **write(A)** instruction of **T2** in schedule 3 does not conflict with the **read(B)** instruction of **T1**, we can swap these instructions to generate an equivalent schedule

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	
	read(B) write(B)

Schedule 3: Only read and write operations

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B)	
write(B)	read(B) write(B)

Schedule 4: After swapping

Conflict Serializability

- Schedules 3 and 5 both produce the same final system state.
- We continue to swap non conflicting instructions:
 - ✓ **Swap the read(B) instruction of T1 with the read(A) instruction of T2.**
 - ✓ **Swap the write(B) instruction of T1 with the write(A) instruction of T2**
 - ✓ **Swap the write(B) instruction of T1 with the read(A) instruction of T2**

T_1	T_2
read(A)	
write(A)	
	read(A)
read(B)	
	write(A)
write(B)	
	read(B)
	write(B)

Schedule 5

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

Schedule 6: Final Schedule

Conflict Serializability

S	
T1	T2
R(A)	
W(A)	
	R(A)
W(A)	
R(B)	

CONFLICT

No change in
Position of Instruction

$S \xrightarrow{CE} S' \rightarrow \text{SERIALIZABLE}$

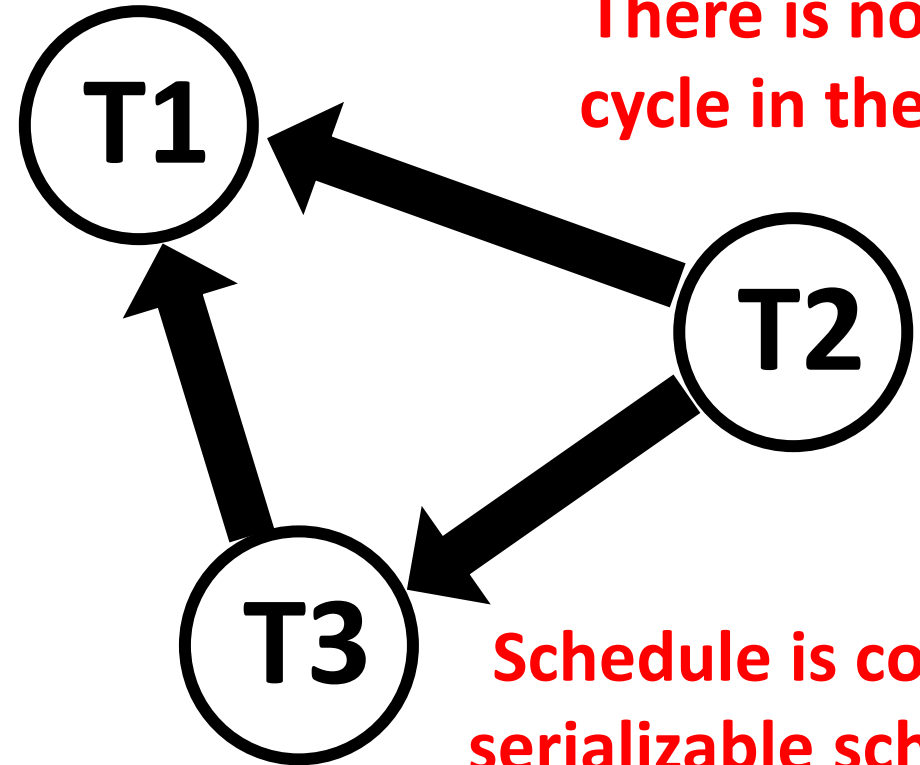
Conflict Serializability

	T1	T2	T3
R – W	R(x)		R(y) R(x)
W – R		R(y) R(z)	
W – R		W(z)	W(y)
	R(z) W(x) W(z)		

Precedence Graph

Loop/Cycle exist?

There is no loop/
cycle in the graph



Schedule is conflict
serializable schedule

Check conflict pairs in other
transaction and draw edge

Conflict Serializability