

Concurrency Control

QUERY PROCESSING

➤ When several transactions execute concurrently in the database, the isolation property may no longer be preserved.

➤ The system must control the interaction among the concurrent transactions, this control is achieved through one of a variety of mechanisms called

Concurrency Control Schemes

➤ There are a variety of concurrency-control schemes.

➤ No one scheme is clearly the best; each one has advantages.

➤ In practice, the most frequently used schemes are ***two-phase locking*** and ***snapshot isolation***.

LOCK BASED PROTOCOLS

- In this case data items should be accessed in **mutually exclusive manner**
- **While one transaction is accessing a data item, no other transaction can modify that data item**
- A transaction can access a data item only if it is currently holding a **lock** on that item.

LOCKS

There are various modes in which a data item may be locked, we restrict our attention to two modes:

1. Shared: If a transaction T_i has obtained a **shared-mode lock** (denoted by **S**) on item Q , then T_i can read, but cannot write, Q

2. Exclusive: If a transaction T_i has obtained an **exclusive-mode lock** (denoted by **X**) on item Q , then **T_i can both read and write Q**

- Every transaction has to **request** a lock to the concurrency-control manager in an appropriate mode on data item Q

- The transaction can proceed only after the concurrency-control manager **grants** the lock to the transaction.

COMPATIBILITY FUNCTION

Let A and B represent arbitrary lock modes. Suppose that a transaction T_i requests a lock of mode A on item Q on which transaction T_j ($T_i \neq T_j$) currently holds a lock of mode B . If transaction T_i can be granted a lock on Q immediately, in spite of the presence of the mode B lock, then we say mode A is **compatible** with mode B .

	S	X
S	true	false
X	false	false

Lock – Compatibility Matrix

An element $\text{comp}(A, B)$ of the matrix has the value *true* if and only if mode A is compatible with mode B .

LOCK REQUEST & GRANTING

- A transaction requests a shared lock on data item Q by executing the **lock-S(Q)** instruction. Similarly, a transaction requests an exclusive lock through the **lock-X(Q)** instruction. A transaction can unlock a data item Q by the **unlock(Q)** instruction.
- Consider the banking example. Let A and B be two accounts that are accessed by transactions T_1 and T_2 . Transaction T_1 transfers \$50 from account B to account A (Figure 1). Transaction T_2 displays the total amount of money in accounts A and B ($A + B$) as shown in Figure 2.

LOCK REQUEST & GRANTING

T1:

```
lock-X(B);  
read(B);  
B := B - 50;  
write(B);  
unlock(B);  
lock-X(A);  
read(A);  
A := A + 50;  
write(A);  
unlock(A)
```

Figure 1 Transaction T1.

T2:

```
lock-S(A);  
read(A);  
unlock(A);  
lock-S(B);  
read(B);  
unlock(B);  
display(A + B)
```

Figure 2 Transaction T2

Suppose that the values of accounts *A* and *B* are \$100 and \$200, respectively. If these two transactions are executed serially, either in the order *T1*, *T2* or the order *T2*, *T1*, then transaction *T2* will display the value \$300. If, however, these transactions are executed concurrently, then schedule 1, in Figure 3, is possible.

LOCK REQUEST & GRANTING

T_1	T_2	concurrency-control manager
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T_2)
	unlock(A)	
	lock-S(B)	
		grant-S(B, T_2)
	read(B)	
	unlock(B)	
	display(A + B)	
lock-X(A)		
read(A)		grant-X(A, T_1)
$A := A - 50$		
write(A)		
unlock(A)		

Figure 3 Schedule I

In this case, transaction T_2 displays \$250, which is incorrect. The reason for this mistake is that the transaction **T_1 unlocked data item B too early**, as a result of which T_2 saw an inconsistent state.

LOCK REQUEST & GRANTING

Suppose now that unlocking is delayed to the end of the transaction. Transaction T_3 corresponds to T_1 with unlocking delayed (Figure 4). Transaction T_4 corresponds to T_2 with unlocking delayed (Figure 5).

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A);

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).

Figure 5 Transaction T_4 (transaction T_2 with unlocking delayed)

Figure 4 Transaction T_3 (transaction T_1 with unlocking delayed).

LOCK REQUEST & GRANTING

➤ The locking can lead to an undesirable situation. Consider the partial schedule of Figure 6 for T_3 and T_4 .

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
lock-X(A)	lock-S(A) read(A) lock-S(B)

Figure 6 Schedule 2

- Neither of these transactions can ever proceed with its normal execution. This situation is called **deadlock**.
- When deadlock occurs, the system must roll back one of the two transactions. Once a transaction has been rolled back, the data items that were locked by that transaction are unlocked. These data items are then available to the other transaction, which can continue with its execution.

STARVATION

Suppose a transaction T_2 has a **shared-mode** lock on a data item, and another transaction T_1 requests an **exclusive-mode** lock on the data item. Clearly, T_1 has to wait for T_2 to release the shared-mode lock. Meanwhile, a transaction T_3 may request a **shared-mode** lock on the same data item. The lock request is compatible with the lock granted to T_2 , so T_3 may be granted the shared-mode lock. At this point T_2 may release the lock, but still T_1 has to wait for T_3 to finish. But again, there may be a new transaction T_4 that requests a shared-mode lock on the same data item, and is granted the lock before T_3 releases it. In fact, it is possible that there is a sequence of transactions that each requests a shared-mode lock on the data item, and each transaction releases the lock a short while after it is granted, but T_1 never gets the exclusive-mode lock on the data item. The transaction T_1 may never make progress, and is said to be **starved**.

AVOIDING STARVATION

We can avoid starvation of transactions by granting locks in the following manner: When a transaction T_i requests a lock on a data item Q in a particular mode M , the concurrency-control manager grants the lock provided that:


1. There is no other transaction holding a lock on Q in a mode that conflicts with M .
2. There is no other transaction that is waiting for a lock on Q and that made its lock request before T_i .

Thus, a lock request will never get blocked by a lock request that is made later.

THE TWO-PHASE LOCKING PROTOCOL

One protocol that ensures serializability is the **two-phase locking protocol**. This protocol requires that each transaction issue lock and unlock requests in two phases:

- 1. Growing phase.** A transaction may obtain locks, but may not release any lock.
- 2. Shrinking phase.** A transaction may release locks, but may not obtain any new locks.



Initially, a transaction is in the growing phase. The transaction acquires locks as needed. Once the transaction releases a lock, it enters the shrinking phase, and it can issue no more lock requests.

DEADLOCK HANDLING

➤ Two-phase locking does *not* ensure freedom from deadlock. Observe that transactions $T3$ and $T4$ are two phase, but, in schedule 2 (Figure 6), they are **deadlocked**.

➤ A system is in a **deadlock** state if there exists a set of transactions such that every transaction in the set is waiting for another transaction in the set. More precisely, there exists a set of waiting transactions $\{T_0, T_1, \dots, T_n\}$ such that T_0 is waiting for a data item that T_1 holds, and T_1 is waiting for a data item that T_2 holds, and \dots , and T_{n-1} is waiting for a data item that T_n holds, and T_n is waiting for a data item that T_0 holds. None of the transactions can make progress in such a situation.

DEADLOCK HANDLING

- The only remedy to deadlock is rolling back some of the transactions involved in the deadlock. Rollback of a transaction may be **partial** (a transaction may be rolled back to the point where it obtained a lock whose release resolves the deadlock)
- There are two principal methods for dealing with the deadlock problem:
 1. We can use a **deadlock prevention** protocol to ensure that the system will *never* enter a deadlock state.
 2. We can allow the system to enter a deadlock state, and then try to recover by using a **deadlock detection** and **deadlock recovery** scheme

DEADLOCK PREVENTION

First Approach:

➤ The simplest scheme requires that each transaction locks all its data items before it begins execution. Disadvantages:

(1) hard to predict, before the transaction begins, what data items need to be locked (2) many of the data items may be locked but unused for a long time.

➤ Other Scheme: to impose an ordering of all data items, and to require that a transaction lock data items only in a sequence consistent with the ordering. A variation of this approach is to use a total order of data items, in conjunction with two-phase locking. Once a transaction has locked a particular item, it cannot request locks on items that precede that item in the ordering. This scheme is easy to implement, as long as the set of data items accessed by a transaction is known when the transaction starts execution.

DEADLOCK PREVENTION

Second Approach:

- Uses **preemption** and transaction **rollbacks**.
- To control the preemption, we assign a unique timestamp, based on a counter or on the system clock, to each transaction when it begins.
- The system uses these timestamps only to decide whether a transaction should wait or roll back.
- Locking is still used for concurrency control.
- If a transaction is rolled back, it retains its *old* timestamp when restarted.
- Two different deadlock-prevention schemes using timestamps have been proposed:

DEADLOCK PREVENTION

I. The **wait-die** scheme is a non preemptive technique. When transaction T_i requests a data item currently held by T_j , T_i is allowed to wait only if it has a timestamp smaller than that of T_j (that is, T_i is older than T_j). Otherwise, T_i is rolled back (dies).

Example:

Suppose that transactions T_{14} , T_{15} , and T_{16} have timestamps 5, 10, and 15, respectively. If T_{14} requests a data item held by T_{15} , then T_{14} will wait. If T_{16} requests a data item held by T_{15} , then T_{16} will be rolled back.

DEADLOCK PREVENTION

2. The wound–wait scheme is a preemptive technique. It is a counterpart to the wait–die scheme. When transaction T_i requests a data item currently held by T_j , T_i is allowed to wait only if it has a timestamp larger than that of T_j (that is, T_i is younger than T_j). Otherwise, T_j is rolled back (T_j is *wounded* by T_i)

Returning to our example, with transactions T_{14} , T_{15} , and T_{16} , if T_{14} requests a data item held by T_{15} , then the data item will be preempted from T_{15} , and T_{15} will be rolled back. If T_{16} requests a data item held by T_{15} , then T_{16} will wait

DEADLOCK PREVENTION

Another simple approach is based on lock timeouts:

➤ A transaction that has requested a lock waits for at most a specified amount of time. If the lock has not been granted within that time, the transaction is said to time out, and it rolls itself back and restarts. If there was in fact a deadlock, one or more transactions involved in the deadlock will time out and roll back, allowing the others to proceed.

➤ Easy to implement, and works well if transactions are short and if long waits are likely to be due to deadlocks.

Issues:

✓ Too long a wait results in unnecessary delays once a deadlock has occurred.


✓ Too short a wait results in transaction rollback even when there is no deadlock, leading to wasted resources.

DEADLOCK DETECTION


- Deadlocks can be described precisely in terms of a directed graph called a wait-for graph.
- This graph consists of a pair $G = (V, E)$, where V is a set of vertices and E is a set of edges. The set of vertices consists of all the transactions in the system.
- Each element in the set E of edges is an ordered pair $T_i \rightarrow T_j$. If $T_i \rightarrow T_j$ is in E , then there is a directed edge from transaction T_i to T_j , implying that transaction T_i is waiting for transaction T_j to release a data item that it needs. This edge is inserted in the wait-for graph.
- This edge is removed only when transaction T_j is no longer holding a data item needed by transaction T_i

DEADLOCK DETECTION

➤ A deadlock exists in the system if and only if the wait-for graph contains a cycle. Each transaction involved in the cycle is said to be deadlocked



Wait-For graph with no cycle



Wait-For graph with cycle

➤ When should we invoke the detection algorithm? The answer depends on two factors:

1. How often does a deadlock occur?
2. How many transactions will be affected by the deadlock?

If deadlocks occur frequently, then the detection algorithm should be invoked more frequently. In the worst case, we would invoke the detection algorithm every time a request for allocation could not be granted immediately.

DEADLOCK RECOVERY

➤ When a detection algorithm determines that a deadlock exists, the system must **recover** from the deadlock. The most common solution is to roll back one or more transactions to break the deadlock. Three actions need to be taken:

I. Selection of a victim:

- ✓ Given a set of deadlocked transactions, we must determine which transaction (or transactions) to roll back to break the deadlock.
- ✓ We should roll back those transactions that will incur the minimum cost. Many factors may determine the cost of a rollback, including:
 - a. How long the transaction has computed, and how much longer the transaction will compute before it completes its designated task.
 - b. How many data items the transaction has used.

DEADLOCK RECOVERY

- c. How many more data items the transaction needs for it to complete.
- d. How many transactions will be involved in the rollback.

2. Rollback:

- ✓ How far the transaction should be rolled back. The simplest solution is a total rollback: Abort the transaction and then restart it.
- ✓ However, it is more effective to roll back the transaction only as far as necessary to break the deadlock. Such **partial rollback** requires the system to maintain additional information about the state of all the running transactions. Specifically, the sequence of lock **requests/grants** and **updates** performed by the transaction needs to be recorded. **The selected transaction must be rolled back to the point where it obtained the first of these locks, undoing all actions it took after that point.** The recovery mechanism must be capable of performing such partial rollbacks.

DEADLOCK RECOVERY

3. Starvation.

- ✓ In a system where the selection of victims is based primarily on cost factors, it may happen that the same transaction is always picked as a victim.
- ✓ As a result, this transaction never completes its designated task, thus there is starvation.
- ✓ We must ensure that a transaction can be picked as a victim only a (small) finite number of times.

TIME STAMP BASED PROTOCOLS

With each transaction T_i in the system, we associate a unique fixed timestamp, denoted by $TS(T_i)$. This timestamp is assigned by the database system before the transaction T_i starts execution. If a transaction T_i has been assigned timestamp $TS(T_i)$, and a new transaction T_j enters the system, then $TS(T_i) < TS(T_j)$. There are two simple methods for implementing this scheme:

1. Use the value of the **system clock** as the timestamp; that is, a transaction's timestamp is equal to the value of the clock when the transaction enters the system. (i.e. T_1 comes at time 2:00:00 and T_2 comes at 2:00:25. This arrival time is acting as a timestamp.)
2. Use a **logical counter** that is incremented after a new timestamp has been assigned; that is, a transaction's timestamp is equal to the value of the counter when the transaction enters the system. (i.e. Counter value can be one from 1,2,3,4,5,...,n)

TIME STAMP BASED PROTOCOLS

➤ To implement this scheme, we associate with each data item Q two timestamp values:

- **W-timestamp(Q)** : denotes the largest timestamp of any transaction that executed $\text{write}(Q)$ successfully.
- **R-timestamp(Q)**: denotes the largest timestamp of any transaction that executed $\text{read}(Q)$ successfully.

These timestamps are updated whenever a new $\text{read}(Q)$ or $\text{write}(Q)$ instruction is executed.

TIME STAMP BASED PROTOCOLS

The Timestamp-Ordering Protocol:

The **timestamp-ordering protocol** ensures that any conflicting read and write operations are executed in timestamp order. This protocol operates as follows:

1. Suppose that transaction T_i issues $\text{read}(Q)$

- a. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten. Hence, the read operation is rejected, and T_i is rolled back.
- b. If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the read operation is executed, and $R\text{-timestamp}(Q)$ is set to the maximum of $R\text{-timestamp}(Q)$ and $TS(T_i)$.

TIME STAMP BASED PROTOCOLS

2. Suppose that transaction T_i issues $\text{write}(Q)$.

- a. If $\text{TS}(T_i) < \text{R-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced. Hence, the system rejects the write operation and rolls T_i back.
- b. If $\text{TS}(T_i) < \text{W-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q . Hence, the system rejects this write operation and rolls T_i back.
- c. Otherwise, the system executes the write operation and sets $\text{W-timestamp}(Q)$ to $\text{TS}(T_i)$.

If a transaction T_i is rolled back by the concurrency-control scheme as result of issuance of either a read or write operation, the system assigns it a new timestamp and restarts it.

TIME STAMP BASED PROTOCOLS

T1	T2	T3
100	200	300
R(A)	R(B) W(B)	R(B) W(A)
W(C)		
R(C)		

	A	B	C
RTS	0	0	0
WTS	0	0	0

1. Suppose that transaction T_i issues read(Q)

a. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten. Hence, the read operation is rejected, and T_i is rolled back.

b. If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the read operation is executed, and $R\text{-timestamp}(Q)$ is set to the maximum of $R\text{-timestamp}(Q)$ and $TS(T_i)$.

2. Suppose that transaction T_i issues write(Q)

a. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced. Hence, the system rejects the write operation and rolls T_i back.

b. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q . Hence, the system rejects this write operation and rolls T_i back.

c. Otherwise, the system executes the write operation and sets $W\text{-timestamp}(Q)$ to $TS(T_i)$.

VALIDATION BASED PROTOCOLS

- In **optimistic concurrency control techniques**, also known as **validation** or **certification techniques**, no checking is done while the transaction is executing.
- In this scheme, updates in the transaction are not applied directly to the database items until the transaction reaches its end.
- During transaction execution, **all updates are applied to local copies of the data items that are kept for the transaction.**
- At the end of transaction execution, a validation phase checks whether any of the transaction's updates violate serializability **If serializability is not violated, the transaction is committed and the database is updated from the local copies; otherwise, the transaction is aborted and then restarted later.**

VALIDATION BASED PROTOCOLS

There are three phases for this concurrency control protocol:

1. Read phase: A transaction can read values of committed data items from the database. However, updates are applied only to local copies (versions) of the data items kept in the transaction workspace.

2. Validation phase: Checking is performed to ensure that serializability will not be violated if the transaction updates are applied to the database.

3. Write phase: If the validation phase is successful, the transaction updates are applied to the database; otherwise, the updates are discarded and the transaction is restarted

VALIDATION BASED PROTOCOLS

Three different timestamps are associated with each transaction T_i :

1. **Start**(T_i), the time when T_i started its execution.
2. **Validation**(T_i), the time when T_i finished its read phase and started its validation phase.
3. **Finish**(T_i), the time when T_i finished its write phase.

We determine the serializability order by the timestamp-ordering technique, using the value of the timestamp $\text{Validation}(T_i)$. Thus, the value $\text{TS}(T_i) = \text{Validation}(T_i)$

VALIDATION BASED PROTOCOLS

The **validation test** for transaction T_i requires that, for all transactions T_k with $TS(T_k) < TS(T_i)$, **one of the following two conditions must hold**:

1. $Finish(T_k) < Start(T_i)$. Since T_k completes its execution before T_i started, the serializability order is indeed maintained.

2. The set of data items written by T_k does not intersect with the set of data items read by T_i and T_k completes its write phase before T_i starts its validation phase ($Start(T_i) < Finish(T_k) < Validation(T_i)$).

This condition ensures that the writes of T_k and T_i do not overlap. Since the writes of T_k do not affect the read of T_i , and since T_i cannot affect the read of T_k , the serializability order is indeed maintained.

VALIDATION BASED PROTOCOLS

T25:

read(*B*);

read(*A*);

display(*A* + *B*)

T26:

read(*B*);

B := *B* − 50;

write(*B*);

read(*A*);

A := *A* + 50;

write(*A*);

display(*A* + *B*)

Transaction *T26* transfers \$50 from account *B* to account *A*, and then displays the contents of both:

VALIDATION BASED PROTOCOLS

Suppose that $TS(T_{25}) < TS(T_{26})$

T_{25}	T_{26}
read(B) read(A) < validate > display($A + B$)	read(B) $B := B - 50$ read(A) $A := A + 50$ < validate > write(B) write(A)

Note: The writes to the actual variables are performed only after the validation phase of T_{26} . Thus, T_{25} reads the old values of B and A , and this schedule is serializable.